



NYMBL · RFQ RESPONSE · 2026

A&M NACR

Restructuring Platform · Phase 1 Development Partner. Prepared for Alvarez & Marsal's North America Corporate Restructuring practice. This response covers our point of view on your eight questions, all seven required response items, and the shape of delivery: a 4-week Discovery + Build, a fixed-fee Engage production rollout, and a delivery pod shipping priced modules — cost and benefit at every gate.

PREPARED FOR

Alvarez & Marsal · NACR

PREPARED BY

Nymbbl

STATUS

Highly Confidential

Highly Confidential. Prepared by Nymbbl for Alvarez & Marsal, North America Corporate Restructuring. This document contains proprietary and commercial information and may be reproduced for the purpose of evaluating this response, but is not to be shared with other parties.

CONTENTS




What's inside.

01	Ingestion & the Semantic Layer	3
02	How This Response Maps to Your RFQ	5
03	Point of View · Q1–Q8	6
04	The Shape of the Engagement	13
05	4-Week Discovery + Build	14
06	Phase-1 Delivery Gates	16
07	Proposed Team	18
08	Engagement Approach & Budget	19
09	Module Scoping & Pricing	21
10	Data Handling · IP · Integrations	23
11	Firm Overview	24
12	Relevant Experience	25

This document follows the presentation: the ingestion thesis, the RFQ mapping, direct answers to all eight questions, the shape of delivery with cost and benefit at every gate, and the team, data handling, and firm credentials behind it.

Get the data in and usable, and everything unlocks.

A&M's own framing is the right one. From the moment a file arrives, it follows one path: land and normalize, map to the semantic layer under human control, then split into two homes.



The crux is the mapping step: an AI first pass, corrected and approved by a practitioner, then frozen as a deterministic rule per client and ERP. The next page details each step.

From landing to two homes.

1 · Land anything, normalize the form. Whatever arrives (Excel, CSV, PDF, Word, PowerPoint) is landed immutably, then normalized: structured files are parsed, including multi-sheet, pivoted, and cryptic-header layouts, and documents run through document AI and OCR to extract fields. This is the same document-ingestion pipeline we are already building for A&M's contract ingestion, and the shape we run in production for a regulated advisory firm's document-management system: a governed repository with sensitivity and retention controls, SharePoint/Graph sync, and AI extraction.

2 · Map to a semantic layer, the crux. Every client's files differ: different ERPs, column names, definitions. We define A&M's semantic layer, the "golden" data model stating what each canonical field means, its type, and whether AI may read it, then have AI take a first pass at mapping a new client's columns using only headers and a few sample rows. A practitioner reviews and corrects; those corrections train the mapping so the next client lands closer to fully mapped. The probabilistic step happens once, under human control, then freezes into a deterministic rule that reproduces identical output on every re-load. Realistically ~70% is common across clients and maps automatically; the remaining ~30% is handled with custom fields and light human review, exactly the WeWork-scale variability A&M described.

3 · Two homes for what comes out. Structured, reportable data lands in the relational data layer (PostgreSQL) that analytics and motion calculations run against. Curated, reusable knowledge (summaries, precedent, mapping crosswalks, methodologies) is derived, never the raw files, and stored in the Knowledge Base as governed, Git-versioned, AI-gated markdown. That derive-don't-hoard split is precisely what lets A&M keep learnings while purging raw client data at engagement close.

CAPTURED WITH YOUR SMES


The final calculation, validation, and exhibit rules for each motion are captured with your SMEs in Discovery, against real historical filings and scrubbed source data.

How this response maps to your RFQ.

We start from one principle, because it determines every downstream decision:

AI accelerates and structures; deterministic engines compute; humans attest.

Every motion runs the same chain.



■ AI-assisted (always reviewed)
 ■ Deterministic engine (filed figures)
 ■ Human attestation gate

YOUR ASK	WHERE WE ANSWER IT
8 questions (“What We Want to Learn From You”)	Point of View, Q1–Q8
1 • Firm overview	Firm Overview
2 • Relevant experience (2–3 case studies)	Relevant Experience
3 • Point of view + discovery approach	Point of View • Discovery
4 • Proposed team	Proposed Team
5 • Indicative engagement approach	Engagement Approach & Budget
6 • Data-handling philosophy	Data Handling
7 • IP confirmation	IP Confirmation

We answer all eight questions and provide all seven required response items. The Point of View section answers each question directly, in a paragraph or less.

A&M's questions, from the RFQ.

We start from one principle, because it determines every downstream decision:

AI accelerates and structures; deterministic engines compute; humans attest.

QUESTION 1 · DATA INGESTION AND TRANSFORMATION

"How would you approach ingesting and transforming highly variable client financial and operational data (i.e. different ERPs, different account structures, different naming conventions, different data sources and sometimes different languages)?"

We treat ingestion as a staged ELT problem with schema-on-read. Source files (PDF, Excel, CSV, ERP extracts) land as-is and immutably in an isolated project room, extracted per file type by reusable parsers and document AI. Account structures and naming are resolved against a canonical restructuring data model using an AI-assisted crosswalk done once, under human approval, then frozen as a deterministic rule per client and ERP. The probabilistic step happens at setup, and every subsequent re-ingestion (for example, weekly actuals) is reproducible. Multiple languages are translated at extraction, with the original retained and mapped to language-independent canonical concepts. This is the acknowledged hard part, and it is the centerpiece of Phase 1.



REPRODUCIBLE › The probabilistic step happens once, at setup, under human control; weekly actuals and every re-ingestion reproduce identical output against the canonical model

Data architecture.

QUESTION 2 · DATA ARCHITECTURE

"How would you design the data architecture to store the case related data in a way that can drive the downstream analytics and outputs?"

Two distinct layers, deliberately separated. A relational data layer (per-module PostgreSQL, versioned and audited, with a Bronze/Silver/Gold "medallion" model and a provenance graph) holds a case's actual structured data. A separate Knowledge Base (AI-gated markdown plus links to source systems, Git-versioned) holds precedent, templates, crosswalks, validation rules, and summaries. Raw documents sit in an object store; retrieval runs on a per-engagement vector store. Gold marts feed both the deterministic outputs and the natural-language interface, so every figure in a motion is traceable from final exhibit back to the exact source record. The shape is one shared foundation, the "spine," that every module snaps onto, shown below.

INTEGRATIONS Documents (Box · SharePoint · Intralinks · VDRs) | Client systems (ERPs · Bank · Treasury) | Courts (Court systems · PACER) | Market data (Octus · Cap IQ · Bordex) | People data (Payroll · Census)

per module, via gateway

SPINE — one shared foundation, every module snaps on

Identity & Access (Entra ID / SSO · project-room isolation · ethical walls)

Document Ingestion (files in, queryable out, per file type · parsers + document AI)

Relational Data Layer (per-engagement PostgreSQL · Bronze/Silver/Gold · versioned · audited · provenance graph)

Tiered AI Routing (Claude: non-confidential · PII scrubbed or in-tenant Azure OpenAI endpoint, never external · logged)

Human-in-the-Loop Controls (confidence & flags · review & attest · immutable audit)

reads & writes derived knowledge only

KNOWLEDGE BASE AI-gated markdown + links to source systems · templates · crosswalks · validation rules · methodologies · precedent language · parameterized by industry / jurisdiction / entity type · Git-versioned · never raw client data

snap onto the spine

MODULES First-Day Motions (vendors, utilities, ...) · Diligence Q&A (grounded, cited answers) · Engage (client data-intake hub) · BART (case management) · Future modules (the module factory)

Governed practitioner datasets.

QUESTION 3 · GOVERNED PRACTITIONER DATASETS

"Not every analysis can be templated. How would you enable A&M practitioners — not developers — to structure and combine ingested files into custom governed datasets (join, tag, add computed fields), analyze them quickly, and promote what they build into reusable templates for future engagements?"

We give practitioners a governed, drag-and-drop workbench: select ingested files and canonical entities, join them (with AI-suggested keys), tag, filter, and add computed fields in a sandboxed formula language, with no arbitrary code. Every dataset is governed by default: versioned, lineage-tracked, deterministic, and inheriting the case's isolation. Practitioners then promote a useful dataset or analysis (for example, a motion's support-binder tab) into a parameterized, version-controlled template in the Knowledge Base; the next engagement runs it against new data after the crosswalk step. Promoted templates become the firm's reusable asset and drive the ~80% reuse across engagements, and provenance is preserved, so an exhibit built from a practitioner dataset still traces to source.

THE RESULT

Practitioner-built, governed datasets promote into parameterized templates in the Knowledge Base. That is what drives the ~80% reuse across engagements, while every exhibit still traces to source.

Accuracy for court-filed outputs.

QUESTION 4 · ACCURACY FOR COURT-FILED OUTPUTS

"How would you design an AI-assisted analytics workflow for outputs that will be filed with a federal court, where there is zero tolerance for material error? What combination of automation, validation, and human review would you propose, and why?"

We separate computation from generation: filed numbers come only from deterministic, tested calculation engines, never generated by an LLM. AI handles extraction, structuring, mapping suggestions, discrepancy flagging, and narrative drafting, all reviewed. Every figure carries a click-through citation to source, so a reviewer verifies rather than re-keys, and a mandatory human attestation gate (preparer, then reviewer, then MD sign-off) stands before anything is court-ready. Reproducibility is guaranteed: same inputs plus same template version produce identical output, and a golden-dataset regression runs against historical filings.

The agent patterns are concrete and drawn from our A&M work: a **Smart Import Validator** that suggests corrections before ingestion, a **Drafting Assistant** that generates first drafts of documents like SOFAs and MORs, and an **Audit Trail Analyzer** that flags anomalies and reconciliation gaps across cases. Each assists the practitioner and never overrides them. Validation itself is defense in depth: every filed figure passes the five layers below before it is court-ready.

- | | |
|---|---|
| 1 | Ingestion checks
Schema, type, completeness, and reconciliation: trial-balance and subledger-to-GL ties |
| 2 | Motion business rules
Statutory caps, date windows, and eligibility for the specific motion |
| 3 | Cross-source anomaly detection
AI flags discrepancies across AP, contracts, and the GL |
| 4 | Control-total reconciliation
Every relief bucket ties back to independent control totals |
| 5 | Human attestation
Preparer, then reviewer, then MD sign-off; nothing is court-ready without it |

COURT-READY › Reproducible from source, with figure-level provenance and a complete audit trail

Diligence Q&A, grounded.

QUESTION 5 · DILIGENCE Q&A, GROUNDED

"How would you approach a natural-language interface for diligence questions against case data, given that hallucinated or unsupported answers are not acceptable?"

Answers are grounded strictly in the case's own documents and data, engagement-scoped. Quantitative questions route through a text-to-query layer that runs a deterministic query against the canonical model and shows the query and its result; there is no free-form LLM arithmetic. Qualitative questions retrieve passages and answer with mandatory citations; when evidence is absent or ambiguous, the system abstains and points to what is missing rather than inferring. Groundedness and citation quality are continuously evaluated and red-teamed, and anything destined for a filing still passes the human-attestation gate.

NO FREE-FORM ARITHMETIC

Quantitative answers come from deterministic queries against the canonical model; qualitative answers carry mandatory citations, and the system abstains when evidence is missing. Anything destined for a filing still passes the human-attestation gate.

Confidentiality and isolation.

QUESTION 6 · CONFIDENTIALITY AND ISOLATION

"How would you design for data confidentiality and isolation across multiple concurrent, highly sensitive client engagements running on the same platform?"

Each engagement runs in a logically isolated project room at every layer: separate schema or instance, object-storage containers, vector namespaces, and encryption keys via Key Vault, so deletion equals key destruction plus purge. PII and census data sit in a separately controlled section, scrubbed to de-identified unique identifiers and/or processed only through an in-tenant Azure OpenAI endpoint, never sent to Claude or any external foundational model. Access is governed by Entra ID/SSO, RBAC, engagement-scoped entitlements, and ethical walls, with a full immutable audit. The AI router enforces and logs that confidential content stays in-tenant; the confidential-versus-cloud boundary is a policy A&M sets. No engagement's data, or model context, is ever visible to another.

ACCESS › Entra ID/SSO · RBAC · engagement-scoped entitlements · ethical walls · immutable audit

ENGAGEMENT A · PROJECT ROOM

Own schema / instance · relational data layer

Own object storage · raw documents

Own vector namespace · engagement-scoped retrieval

Own encryption keys · Key Vault; deletion = key destruction + purge

PII · separately controlled › scrubbed to unique IDs and/or in-tenant Azure OpenAI · never external

ENGAGEMENT B · PROJECT ROOM

Own schema / instance · relational data layer

Own object storage · raw documents

Own vector namespace · engagement-scoped retrieval

Own encryption keys · Key Vault; deletion = key destruction + purge

PII · separately controlled › scrubbed to unique IDs and/or in-tenant Azure OpenAI · never external

AI ROUTER › Enforces and logs every routing decision · confidential content stays in-tenant · the boundary is a policy A&M sets

Knowledge base and delivery sequencing.

QUESTION 7 · BUILDING THE KNOWLEDGE BASE

"Given that A&M practitioners hold the domain expertise, how would you approach building and evolving a knowledge base that captures and applies that expertise?"

The Knowledge Base codifies your practitioners' expertise as AI-gated, versioned markdown plus links: a template library, mapping crosswalks, validation-rule sets, diligence checklists, calculation methodologies, and precedent language, each parameterized by industry, jurisdiction, and entity type. It is actively gatekept (it rejects duplicate descriptions and blocks a module from authoring a field it does not own) so it never degrades into a data lake, and it references source rather than hoarding raw data. It grows three ways: practitioners promoting datasets and analyses to templates, SMEs authoring rules and checklists, and learning from corrections, where an approved reviewer override becomes a rule or example. Work done in one motion or module compounds into the next; A&M practitioners own the domain input, the platform operationalizes it.

QUESTION 8 · SEQUENCING DELIVERY

"How would you sequence delivery, given the tension between speed (the case for automation is strongest where deadlines are tightest) and the accuracy bar described above?"

We resolve the tension by proving the whole chain end to end on the narrowest scope first. A 4-week Discovery + Build engineers the context and ships a working data-integration workflow for the 13-week cash flow use case on A&M's own files; the Engage 2.0 production rollout (3 months) then takes ingestion plus the data mapper and parser — the client data-intake front end A&M calls "step one" of the whole process — into production, starting with the ingestion A&M itself calls the unlock. Utilities follows as the next motion; vendors and utilities are the two A&M named. Every motion follows the same shape A&M articulated: ingest the data, prepare relief calculations by bucket, each bucket becomes a tab in a support binder, then insert the resulting language, tables, and exhibits into the filing. Build one motion well and motions 2–N are largely configuration; that is the path to ~10 first-day-motion modules. Speed comes from the shared foundation and automation; accuracy comes from validation and human-review gates that never move. We then parallel-path so value ships while the foundation hardens.


Our discovery approach is detailed in the Discovery section: a focused 4-week context-engineering sprint that locks sources, rules, templates, and the security and AI-routing boundary for the vendor and utilities motions, curates the golden-case regression suite, and resolves the open decisions before build accelerates.

04 DELIVERY SCHEDULE



The shape of the engagement.

Three phases, each de-risking the next. Phase 1 — Discovery + Build — is a firm \$60K fixed fee committed now: four weeks that clear build scope and ship a working ingestion → mapping → forecast flow for the 13-week cash flow use case, on A&M's own files. Phase 2 puts Engage 2.0 into production — ingestion plus the data mapper and parser — for a fixed \$190K over three months. Phase 3 stands up the delivery pod: four FTEs shipping priced modules at case speed, \$100K a month with a 12-month minimum — and no multi-year commitment upfront.



COST AND BENEFIT AT EVERY GATE

Phase 1 is committed now at a firm \$60K; each later phase is priced at its gate — cost and benefit, phase by phase, with no multi-year commitment upfront.

05 DISCOVERY



4-Week Discovery + Build.

A compressed four-week sprint — discovery that also builds. It clears build scope, produces the data-handling protocols and firm pricing, and ships a working data-integration workflow for the 13-week cash flow use case on A&M's own files.



Six workstreams across the four weeks, with the Phase-1 delivery gates on the same clock: M1 data in and structured (~week 2), M2 mapping ratified and the forecast ties (~week 3), and M3 the working flow demonstrated (~week 4).

Week 1 · Source & access inventory. Every system of record and access path: ERPs including the cash-disbursements extract, Box, SharePoint, Intralinks, VDRs, PACER, payroll and census, bank feeds, and market data, and how cleanly each connects.

Week 2 · Canonical model v1. The entities the first workflows need, validated with SMEs.

Week 3 · Engage front-end scoping. Confirm Engage as the production data-intake front end, A&M's "step one" for getting client files into one place, and its rebuild-versus-enhance scope. The manual diligence sub-feature is deferred.

Weeks 1-2 · Golden-case curation. Historical engagement data plus filed exhibits and support binders become the build's regression suite and accuracy benchmark. A&M shares first-day-motion examples and scrubbed source data.

Weeks 2-3 · Motion selection & support-binder rule capture. A&M named vendors and utilities as the two first-day motions, vendors as the first module target, because together they span enough data types to shape the architecture. We define each one's exact bucket relief calculations, validations, and exhibit format with SMEs.

Week 4 · Security boundary & AI-routing policy. Isolation model, PII handling, and the routing policy that separates PII and confidential content from cloud and Claude inference.

05 DISCOVERY · OPEN DECISIONS

Discovery closes these open decisions so the build starts on settled ground.

PLATFORM & DATA STORE

- **Cloud choice:** Azure is the default, but not locked; the platform must integrate with BART on AWS.
- **Data-store and warehouse choice:** PostgreSQL is our lead recommendation; Fabric, Snowflake, and Databricks are evaluated against it.
- **Platform naming.**

AI & GOVERNANCE


- **Determinism bar:** how deterministic the AI must be for anything touching filings.
- **Human-in-the-loop level** per output type.
- **Retention taxonomy:** purge versus retain-in-aggregate.

SCOPE & NAVIGATION


- **Engage:** rebuild-versus-enhance scope.
- **Platform navigation and role mapping.**
- **Adjacent capabilities:** where a CRM, deal staffing, and data rooms fit.

Three testable gates inside the fixed fee.

Inside the fixed-fee four-week Phase 1 — Discovery + Build — delivery is measured against three testable gates, each demonstrated live.




M1 · Data in and structured (~week 2). Real A&M source files land, parse per file type, and crosswalk to canonical model v1 under practitioner approval. Done when every value traces back to its exact source cell or document.



M2 · Mapping ratified, forecast computes (~week 3). Vendor resolution and the pre/post-petition split land, the practitioner-approved crosswalk freezes as a deterministic rule, and the 13-week forecast computes — date windows and control totals passing or failing on their own; a figure that fails to tie is flagged, never emitted.

Working flow, demonstrated.



M3 · Working flow, demonstrated (~week 4). The end-to-end flow runs on A&M's own files — ingestion → mapping → 13-week forecast — with click-through provenance on every figure and the attestation gate in place. Done when the demonstrated workflow reproduces the curated golden cases within a documented tolerance, with a recorded human sign-off: the Phase-1 exit.

07 PROPOSED TEAM



A lean, senior pod that scales by track.

A deliberately lean, senior pod for the build, scaling by track after.

■ PLATFORM CONSULTANT · OVER THE TOP



Martyn Mason

Platform Consultant · over the top

Solution and domain coherence across all tracks, so the foundation stays genuinely reusable.

■ PHASE-1 DELIVERY POD



Ruben Carrera

Engagement Lead

Delivery ownership, milestones, and the A&M product-owner interface.



Andros Haggins

Technical Architect

Data and integration architecture; security and isolation validation.

Cloud Architect

TBD

Azure infrastructure, per-engagement isolation, AI routing and deployment.

QA Engineer

TBD

The accuracy and regression harness: golden-case reproduction, rule and provenance tests.

■ A&M · DAY-TO-DAY PRODUCT OWNERS



Andrew Ku

NACR · leading this initiative

Day-to-day product owner, setting priorities and requirements with SME input.



John Bain

Director

Day-to-day product owner alongside Andrew, setting priorities and requirements.

ALSO AT THE TABLE

A NACR Managing Director as executive sponsor and domain authority, and a Technical Solutions Architect for architecture validation. Solution Architects and specialists scale in per track after the build. Point of contact for this response: Ruben Carrera (ruben.carrera@nymbl.app). The named team is available for the Stage-2 discovery sessions and can begin as soon as an agreement is in place.

Discovery, then a phased approach.

A four-week Discovery + Build proves the workflow on your files for a fixed \$60K; a fixed-fee rollout puts Engage 2.0 into production; a standing delivery pod then ships priced modules at case speed. Cost and benefit at every gate — no multi-year commitment upfront.



Phase 1 is committed now at a firm \$60K; each later phase is priced at its gate — cost and benefit, phase by phase, with no multi-year commitment upfront. The indicative structure follows on the next page.

Indicative engagement structure.

Each stage with its scope, duration, and commercial shape, as presented.

STAGE	SCOPE & TEAM	DURATION	COMMERCIAL	INDICATIVE BUDGET
Phase 1 · Discovery + Build	Compressed discovery (6 → 4 weeks) that clears build scope and ships a live ingestion → mapping → forecast flow for the 13-week cash flow use case — on your files in month one, not a slide deck.	4 weeks	Fixed fee	\$60K
Phase 2 · Engage production rollout	Full production build of Engage 2.0 — ingestion plus the data mapper and parser, in production across FTM workflows; calendar/PMO deferred.	3 months	Fixed fee	\$190K
Phase 3 · Delivery pod + modules	Standing team of four FTEs — 2 Solution Architects + 2 support — building platform capability delivered as priced modules, at case speed, without waiting on a dev.	ongoing	Retainer · 12-mo min	\$100K/mo
What the retainer buys	Depending on module size and leverage — net of enablement, run-and-maintain, and the discovery the pod also absorbs.	≈ \$1.2M/yr	Pod capacity	3–5 modules/yr

Indicative budgets as presented; Phase 1 is the committed fixed fee.

How modules are scoped and priced.

A module is one governed dataset plus the models built on it — the 13-week cash flow, or the store-footprint analysis. Modules are priced by tier, and every ratified mapping makes the next similar module cheaper.

Small — Leveraged

**1–2 months ·
High reuse**

~80% common workflow;
ratified mappings apply;
structured data only.

Medium — Configured

**2–3 months ·
Partial reuse**

Meaningful net-new
mapping; several messy
source formats; moderate
governance.

Large — Net-new / Complex

**3–5 months ·
Little to no
reuse**

Structured joined to
unstructured docs;
declaration-grade lineage;
iterative model logic.

WHAT THE RETAINER BUYS

\$100K/mo ≈ \$1.2M a year ≈ 3–5 modules — depending on size and leverage, net of enablement, run-and-maintain, and the discovery the pod also absorbs.

09 MODULE PRICING · WHAT SIZES A MODULE



What sizes a module.

Five levers move a module up or down a tier; the sizing table below is those levers applied to the three modules discussed. Directional, not quotes.

THE FIVE LEVERS what moves a module up or down a tier

1	Reuse — the biggest lever New data through an already-ratified mapping ("build once") — each module makes the next similar one cheaper.
2	Source variability Clean keyed extracts vs. free-text payees, multi-bank PDF/CSV, hand-kept logs, duplicate vendor masters.
3	Structured joined to unstructured Joining a GL to lease documents — with lineage to source page and clause — is materially harder.
4	Governance grade Reporting-grade vs. declaration-grade: cell-level lineage and versioning for figures behind a sworn declaration.
5	Enablement depth One-off build vs. codified so a practitioner self-composes it next time.

THE LEVERS, APPLIED illustrative sizing for the three modules discussed

MODULE	WHAT SIZES IT	TIER	INDICATIVE SIZING
13-Week Cash Actuals + Variance	Two joins carry it; vendor resolution, the pre/post-petition split, and payment allocation are the hard parts. Drops toward the low end once Phase 1 ratifies the mapping.	Medium · 2–3 months	~\$200–275K
GL / P&L Financial Analysis	~1M-row reshape, multiple mapping tables, segmentation as dimensions, reconciliation to trial balance.	Large · 3–5 months	~\$325–425K
Store Footprint — financials + leases	Structured joined to unstructured lease terms, source-page/clause lineage, iterative closure-list logic.	Large · 3–5 months	~\$400–500K

Each figure drops as ratified mappings accumulate: the reuse lever working in A&M's favor.

Data handling, ownership, and integrations.

DATA HANDLING PHILOSOPHY

Confidential data stays inside A&M's Azure tenant and is never used to train models. During development and testing we use synthetic or A&M-provided historical (scrubbed) data in an isolated environment; production access is least-privilege and audited. AI routing is tiered by sensitivity: **non-confidential, non-PII inference goes to Claude via A&M's enterprise agreement; PII and confidential content is scrubbed to de-identified unique identifiers and/or processed only through an in-tenant Azure OpenAI endpoint, never sent to Claude or any external foundational model**, with every routing decision logged. The confidential-versus-cloud boundary is a policy A&M sets in Discovery. Retention is treated as a first-class requirement across three tiers: **raw client data and case-derived working data are purged and verifiably deleted at engagement close** (key destruction plus data removal across all stores and backups), while only **de-identified learnings** (precedent, templates, methodologies) persist in the Knowledge Base as A&M's compounding asset. The precise purge-versus-retain boundary is a business decision to finalize in Discovery, where Nymbbl will bring a recommended taxonomy. Nymbbl maintains SOC 2 and aligns its controls to A&M's security requirements.

IP CONFIRMATION

Ownership is how we work by default: we build standard, portable code that belongs to the client, with no proprietary platform to license back. Accordingly, Nymbbl confirms acceptance of A&M's ownership terms **for all custom work product built for A&M under this engagement**: source code, data models, architecture, and documentation created for A&M are A&M's sole property upon creation, and will not be reused for any other client or internal purpose. **"Reusable platform" here means reusable across A&M's own engagements, not portable to other Nymbbl clients.** This does not transfer Nymbbl's **pre-existing IP** (our delivery framework, methods, and generic tooling brought to the engagement), which remains Nymbbl's; A&M receives full rights to use it as embedded in the delivered work product. This confirmation is given subject to legal review of the definitive agreement language.

INTEGRATIONS

Microsoft 365, Teams, SharePoint, and Box are first-class targets alongside client ERPs and bank and treasury feeds; court systems and PACER, market data (Octus, Cap IQ, Bordex), and payroll and census follow by module. MCP is evaluated as a connector pattern. Additional proprietary A&M systems are scoped in Discovery. We connect to the systems A&M already runs rather than rebuilding them; each integration lands per module through the spine's gateway, inheriting the engagement's isolation and audit.

11 FIRM OVERVIEW



Boston-based, AI-native, already in your environment.

Nymbl is a Boston-based (Quincy, MA), AI-native software firm that builds the custom enterprise systems modern organizations run on, for the moment when off-the-shelf almost fits and the workarounds keep piling up. Four capabilities bear directly on this engagement: custom enterprise software development, enterprise AI (agentic and document AI), data readiness, and team augmentation. We are AI-accelerated by default, using AI to compress delivery timelines while keeping the resulting systems production-grade and owned entirely by the client.

We are deliberately anti-transactional, building long-term partnerships across the full lifecycle (strategy, execution, and ongoing growth) rather than one-off projects. Our delivery is strategy-first: we stand up working technical assets early to validate the business strategy rather than deferring engineering to late stages, exactly what this platform's Phase 1 calls for.

We bring twenty years of building enterprise systems, evolved through three eras (data warehousing, low-code application platforms, and now AI-native development), so the acceleration we describe rests on two decades of delivery, not a first bet. Notable clients include Boeing, CACI (as their federal industry application development partner), Houlihan Lokey, Veolia, and Colgate, spanning federal, high-compliance, and financial-advisory contexts directly relevant to this work. We staff deliberately lean per engagement while drawing on a global bench as the platform broadens, and we already deliver inside A&M's NACR environment today (BART).

AT A GLANCE

100+ people across 14 countries · 155+ clients, 20 of them Fortune 500 · SOC 2 maintained · headquartered in the Boston area (Quincy, MA), the office that would lead this engagement.

12 RELEVANT EXPERIENCE

Cleared, on-point proof.

VytalMed

VYTALMED · AI IN PRODUCTION

An AI-native operating system running a live radiology business, in production with zero pilots. Proof we put AI into mission-critical operations, not demonstrations.

THE KRAFT GROUP

THE KRAFT GROUP · REPLATFORMING & OWNERSHIP

Eight legacy applications replatformed in ten months, on time and on budget, eliminating **\$300K/year** in licensing. CIO Michael Israel: "Very rarely do I see a vendor actually deliver everything they said they would, and within budget." Proof of the own-what-we-build, lose-the-licensing model A&M's IP terms call for.



BART · A&M NACR (INCUMBENT)

We already build A&M's case-management platform: claims, schedules and statements of financial affairs, noticing, the authoritative creditor matrix, and court-ready output. Evidence we understand your domain, data, and standards, and already work in your environment.



SELECTED DELIVERY

ADDITIONAL DIRECTLY RELEVANT DELIVERY

Intelligent document processing at scale: 1M+ documents/year with a Customer-360 view exposed to AI agents, the ingestion and OCR backbone this build needs. \$20M in purchase orders validated for a healthcare supply chain, the same class of work as a vendor-payment motion. And a semantic layer turning plain-English questions into governed SQL across chat, Slack, and Excel, the capability behind our answer to Q5.

Phase-1 delivery risk is further reduced because the end-to-end chain (ingest, govern, compute, validate, human review, court-ready) is already proven in a working Nymbbl demo; Phase 1 applies that proven pattern to A&M's real data rather than inventing it.



AI accelerates and structures; deterministic engines
compute; humans attest.

Thank you for the opportunity to respond. The chain this response
describes is not a proposal of research; it is a working demo we would
walk live in the Stage-2 discovery sessions. We are ready to begin as
soon as an agreement is in place.

Ruben Carrera · Engagement Lead · ruben.carrera@nymbl.app